

Comparing Optimization Strategies

Learn From My Mistakes

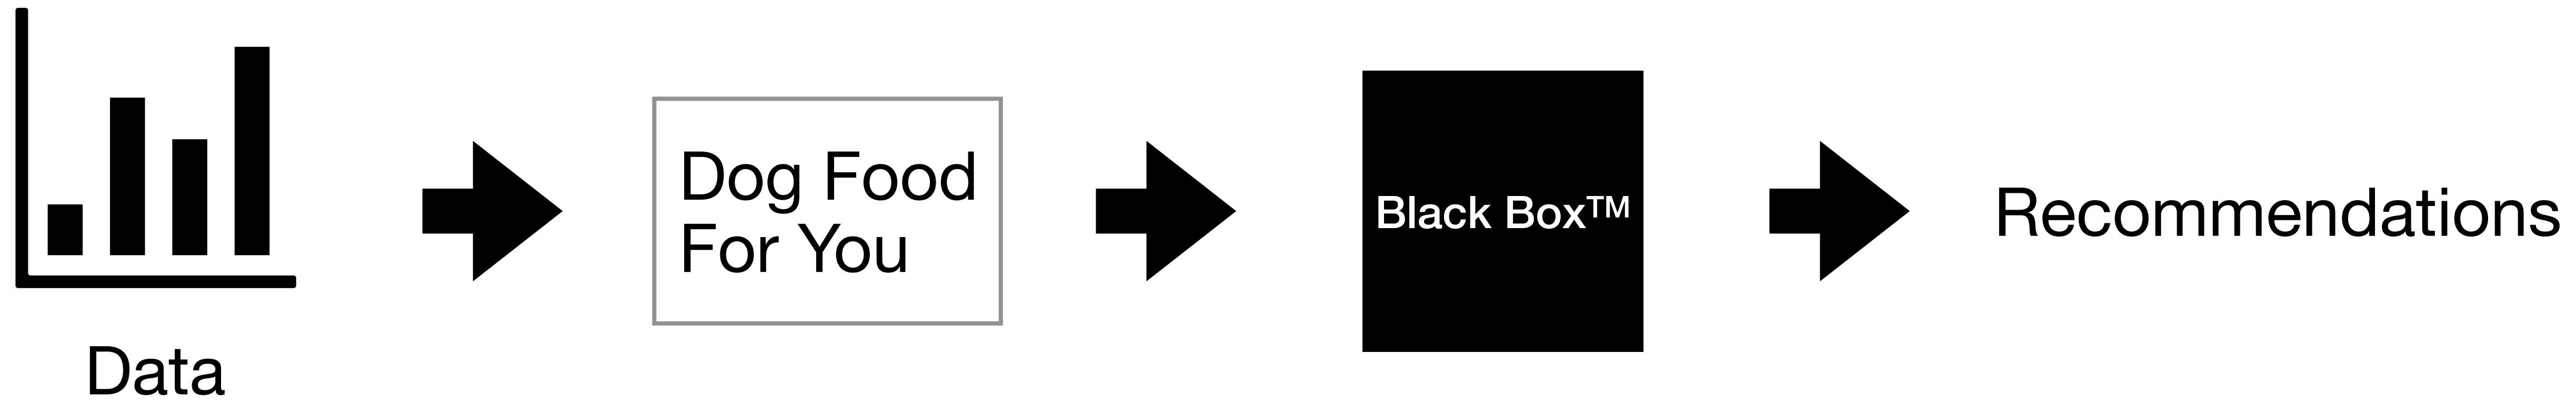
Mark Lodato

Goals

- Talk about optimizing real code
- Compare different approaches taken and their pay-offs
- Show how we can reach for use-case-specific algorithms
- Introduce a handy tool for analyzing code performance
- Start a conversation about optimization

Dog Food For You

State of the Art Black Box™ Technology



Optimization Strategy 1

Guess

```
pub fn write_file(...) {  
    for dog in mapping.dogs() {  
    }  
}
```

Optimization Strategy 1

Guess

```
pub fn write_file(...) {  
  
    for dog in mapping.dogs() {  
  
        for _ in 0..lines_per_dog {  
  
        }  
    }  
}
```

Optimization Strategy 1

Guess

```
pub fn write_file(...) {  
  
    for dog in mapping.dogs() {  
  
        for _ in 0..lines_per_dog {  
  
            let line = walk(...);  
  
            output_file.write_all(...);  
        }  
    }  
}
```

Optimization Strategy 1

Guess

```
pub fn walk(...) -> String {  
    let mut line: Vec<String> = ...;  
  
    // bunch of line.push()s  
  
    line.join(" ") + "\n"  
}
```

Optimization Strategy 1

Guess

<code>test array_concat</code>	... bench:	38 ns/iter (+/- 27)
<code>test array_join</code>	... bench:	38 ns/iter (+/- 22)
<code>test array_join_long</code>	... bench:	35 ns/iter (+/- 18)
<code>test collect_from_array_to_string</code>	... bench:	80 ns/iter (+/- 17)
<code>test collect_from_vec_to_string</code>	... bench:	76 ns/iter (+/- 27)
<code>test format_macro</code>	... bench:	111 ns/iter (+/- 69)
<code>test from_bytes</code>	... bench:	1 ns/iter (+/- 0)
<code>test mut_string_push_str</code>	... bench:	75 ns/iter (+/- 35)
<code>test mut_string_push_string</code>	... bench:	164 ns/iter (+/- 77)
<code>test mut_string_with_capacity_push_str</code>	... bench:	32 ns/iter (+/- 14)
<code>test mut_string_with_capacity_push_str_char</code>	... bench:	29 ns/iter (+/- 12)
<code>test mut_string_with_too_little_capacity_push_str</code>	... bench:	96 ns/iter (+/- 11)
<code>test mut_string_with_too_much_capacity_push_str</code>	... bench:	37 ns/iter (+/- 8)
<code>test string_from_all</code>	... bench:	133 ns/iter (+/- 101)
<code>test string_from_plus_op</code>	... bench:	76 ns/iter (+/- 9)
<code>test to_owned_plus_op</code>	... bench:	81 ns/iter (+/- 15)
<code>test to_string_plus_op</code>	... bench:	83 ns/iter (+/- 17)

Optimization Strategy 1

Guess

test array_concat	... bench:	38 ns/iter (+/- 27)
test array_join	... bench:	38 ns/iter (+/- 22)
test array_join_long	... bench:	35 ns/iter (+/- 18)
test collect_from_array_to_string	... bench:	80 ns/iter (+/- 17)
test collect_from_vec_to_string	... bench:	76 ns/iter (+/- 27)
test format_macro	... bench:	111 ns/iter (+/- 69)
test from_bytes	... bench:	1 ns/iter (+/- 0)
test mut_string_push_str	... bench:	75 ns/iter (+/- 35)
test mut_string_push_string	... bench:	164 ns/iter (+/- 77)
test mut_string_with_capacity_push_str	... bench:	32 ns/iter (+/- 14)
test mut_string_with_capacity_push_str_char	... bench:	29 ns/iter (+/- 12)
test mut_string_with_too_little_capacity_push_str	... bench:	96 ns/iter (+/- 11)
test mut_string_with_too_much_capacity_push_str	... bench:	37 ns/iter (+/- 8)
test string_from_all	... bench:	133 ns/iter (+/- 101)
test string_from_plus_op	... bench:	76 ns/iter (+/- 9)
test to_owned_plus_op	... bench:	81 ns/iter (+/- 15)
test to_string_plus_op	... bench:	83 ns/iter (+/- 17)

Optimization Strategy 1

Guess

```
pub fn walk(..., line: &mut String) {  
  
    // bunch of line.push_str()'s  
    // and line.push(' ')s  
  
    line.push('\'\n');  
}
```

Test it out!

	Run 1 (ms)	Run 2 (ms)	Run 3 (ms)	Run 4 (ms)	Run 5 (ms)
join	59364	58472	59844	58018	58810
push_str	57407	57252	56730	57590	56278
diff	-3.3%	-2.09%	-5.2%	-0.74%	-4.31%



Guessing about performance is hard

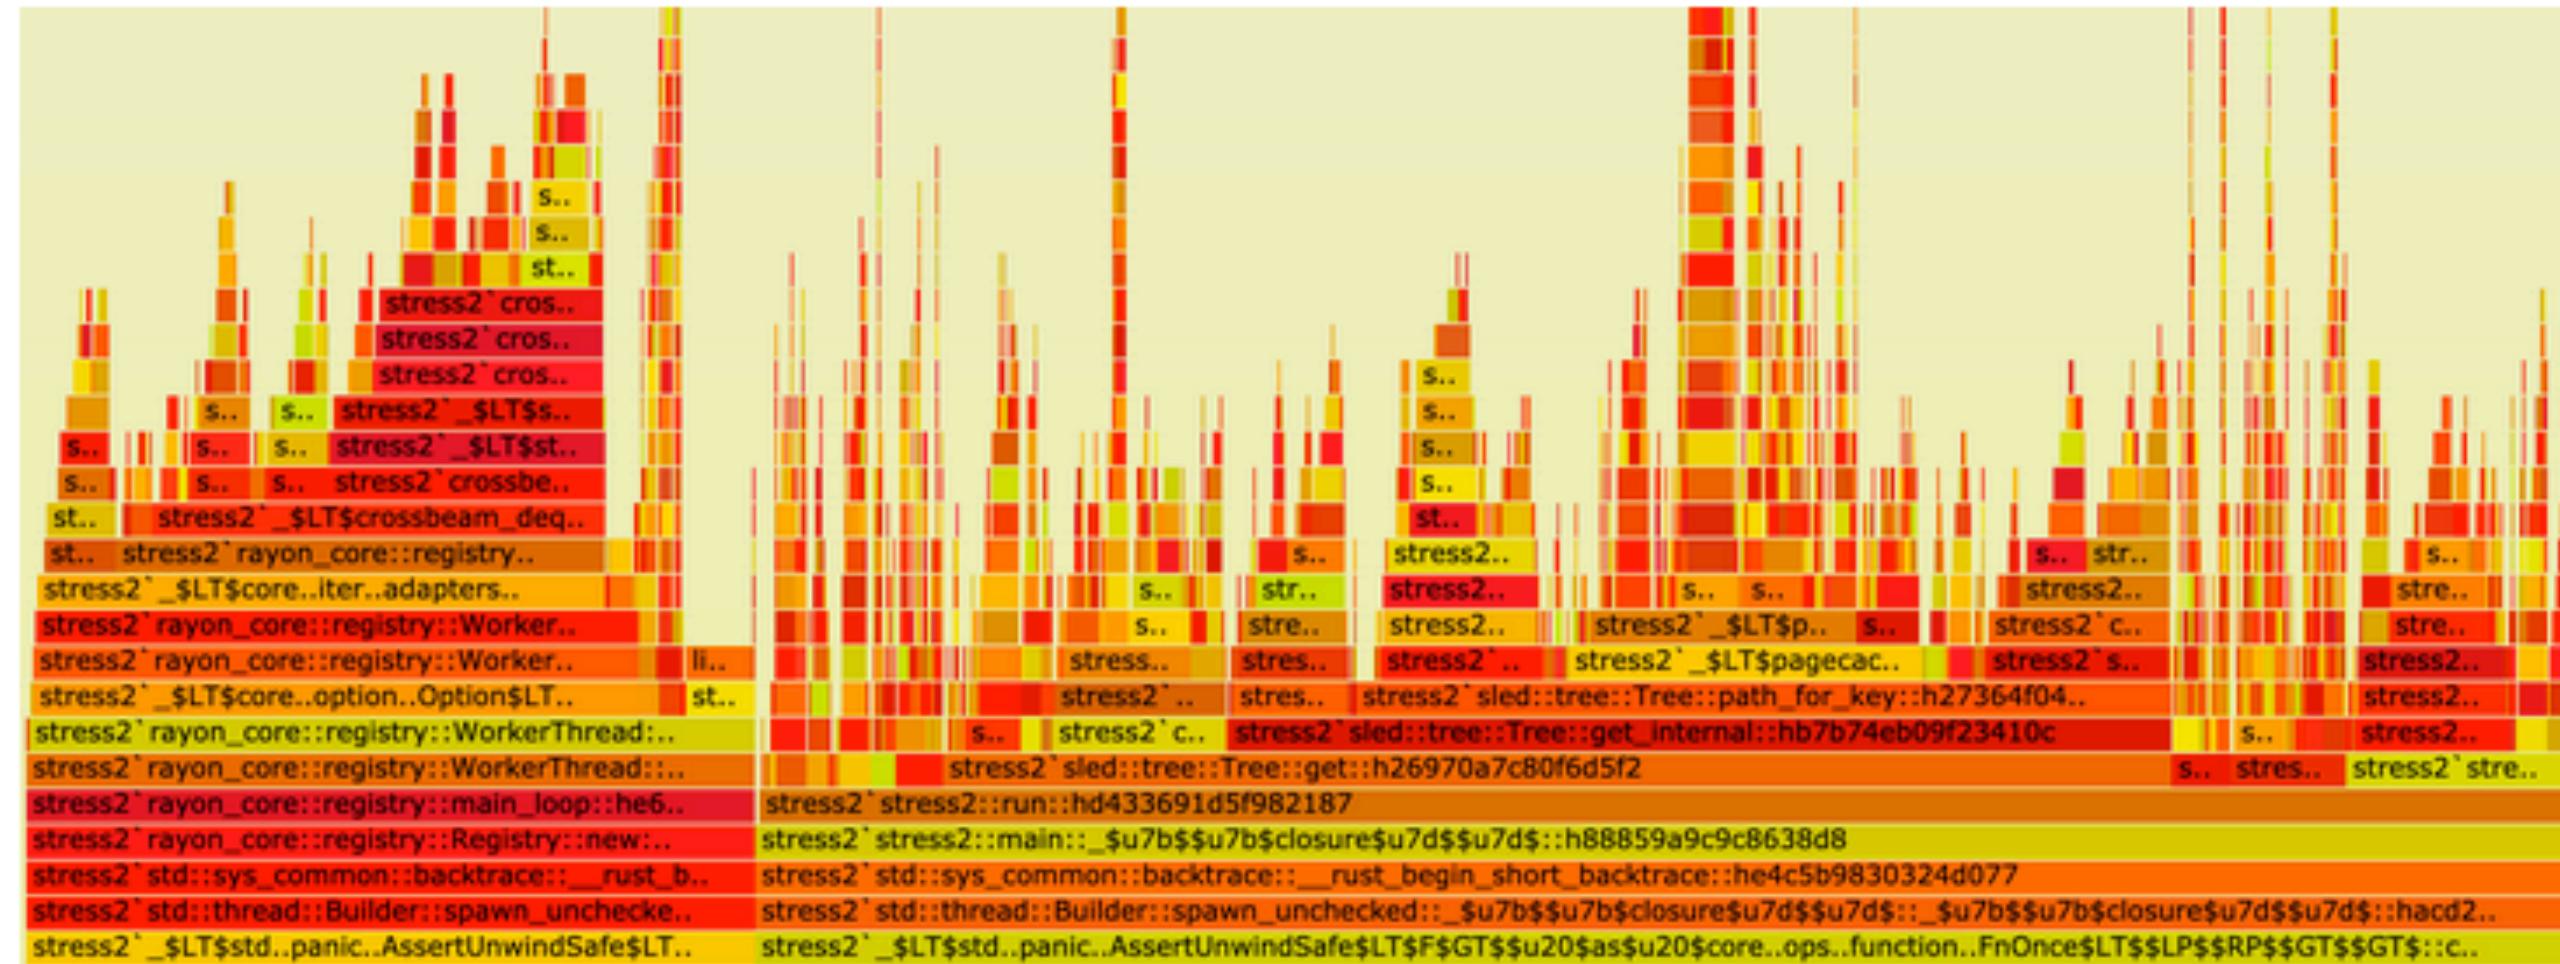
- Higher level languages exist for a reason

Optimization Strategy 2

Measure

README.md

[cargo-]flamegraph



A Rust-powered flamegraph generator with additional support for Cargo projects! It can be used to profile anything, not just Rust projects! No perl or pipes required <3

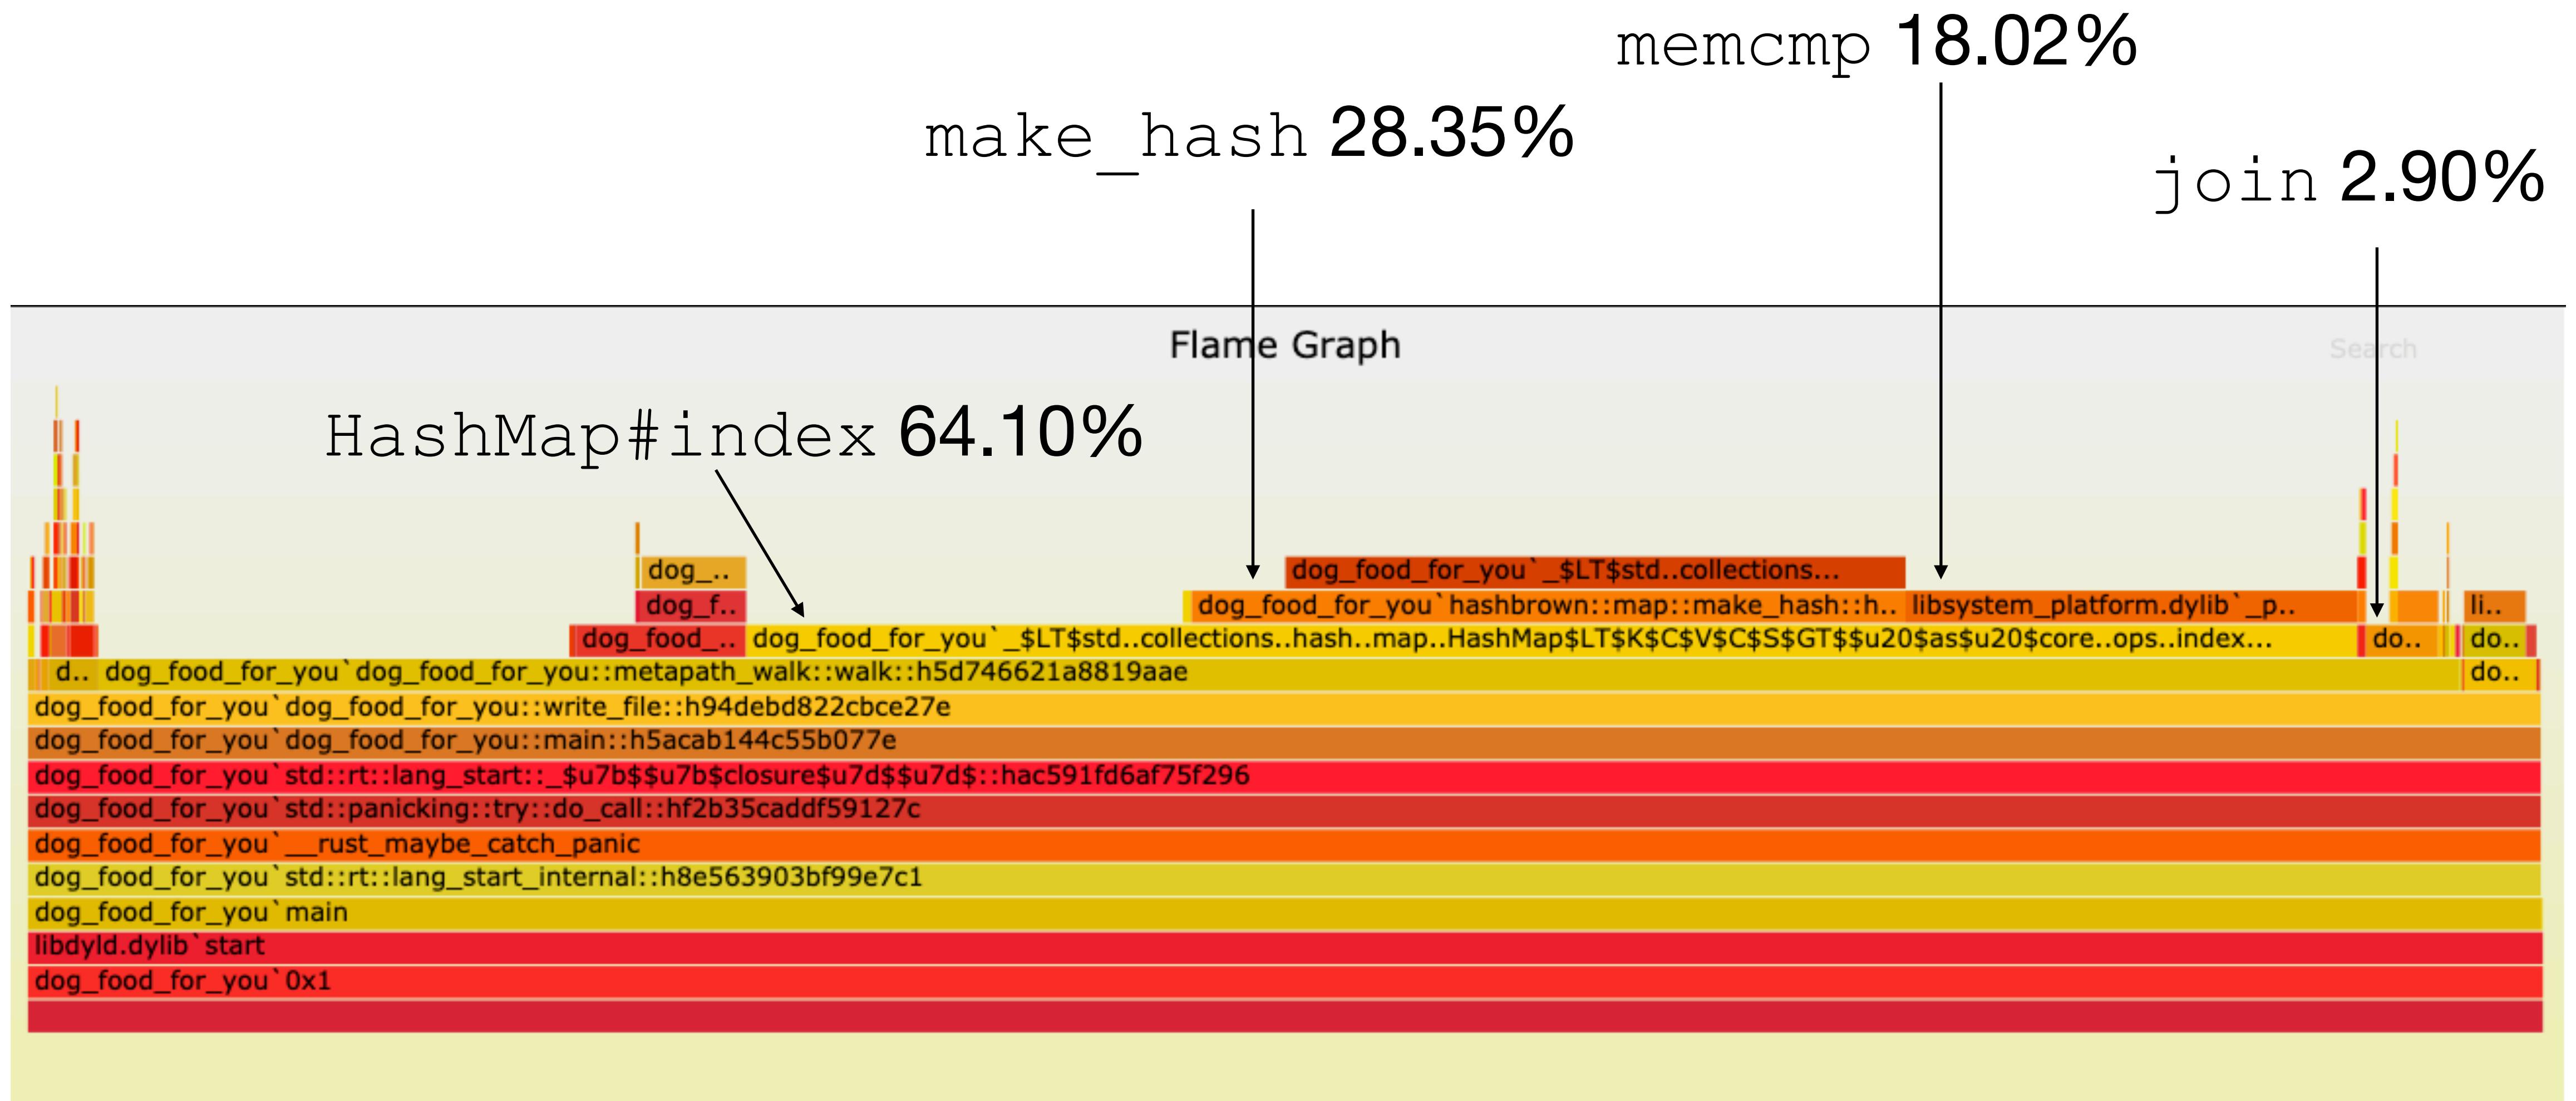
How to use flamegraphs: [what's a flamegraph, and how can I use it to guide systems performance work?](#)

Relies on perf on linux and dtrace otherwise. Built on top of [@jonhoo's](#) wonderful [Inferno](#) all-rust flamegraph generation library!

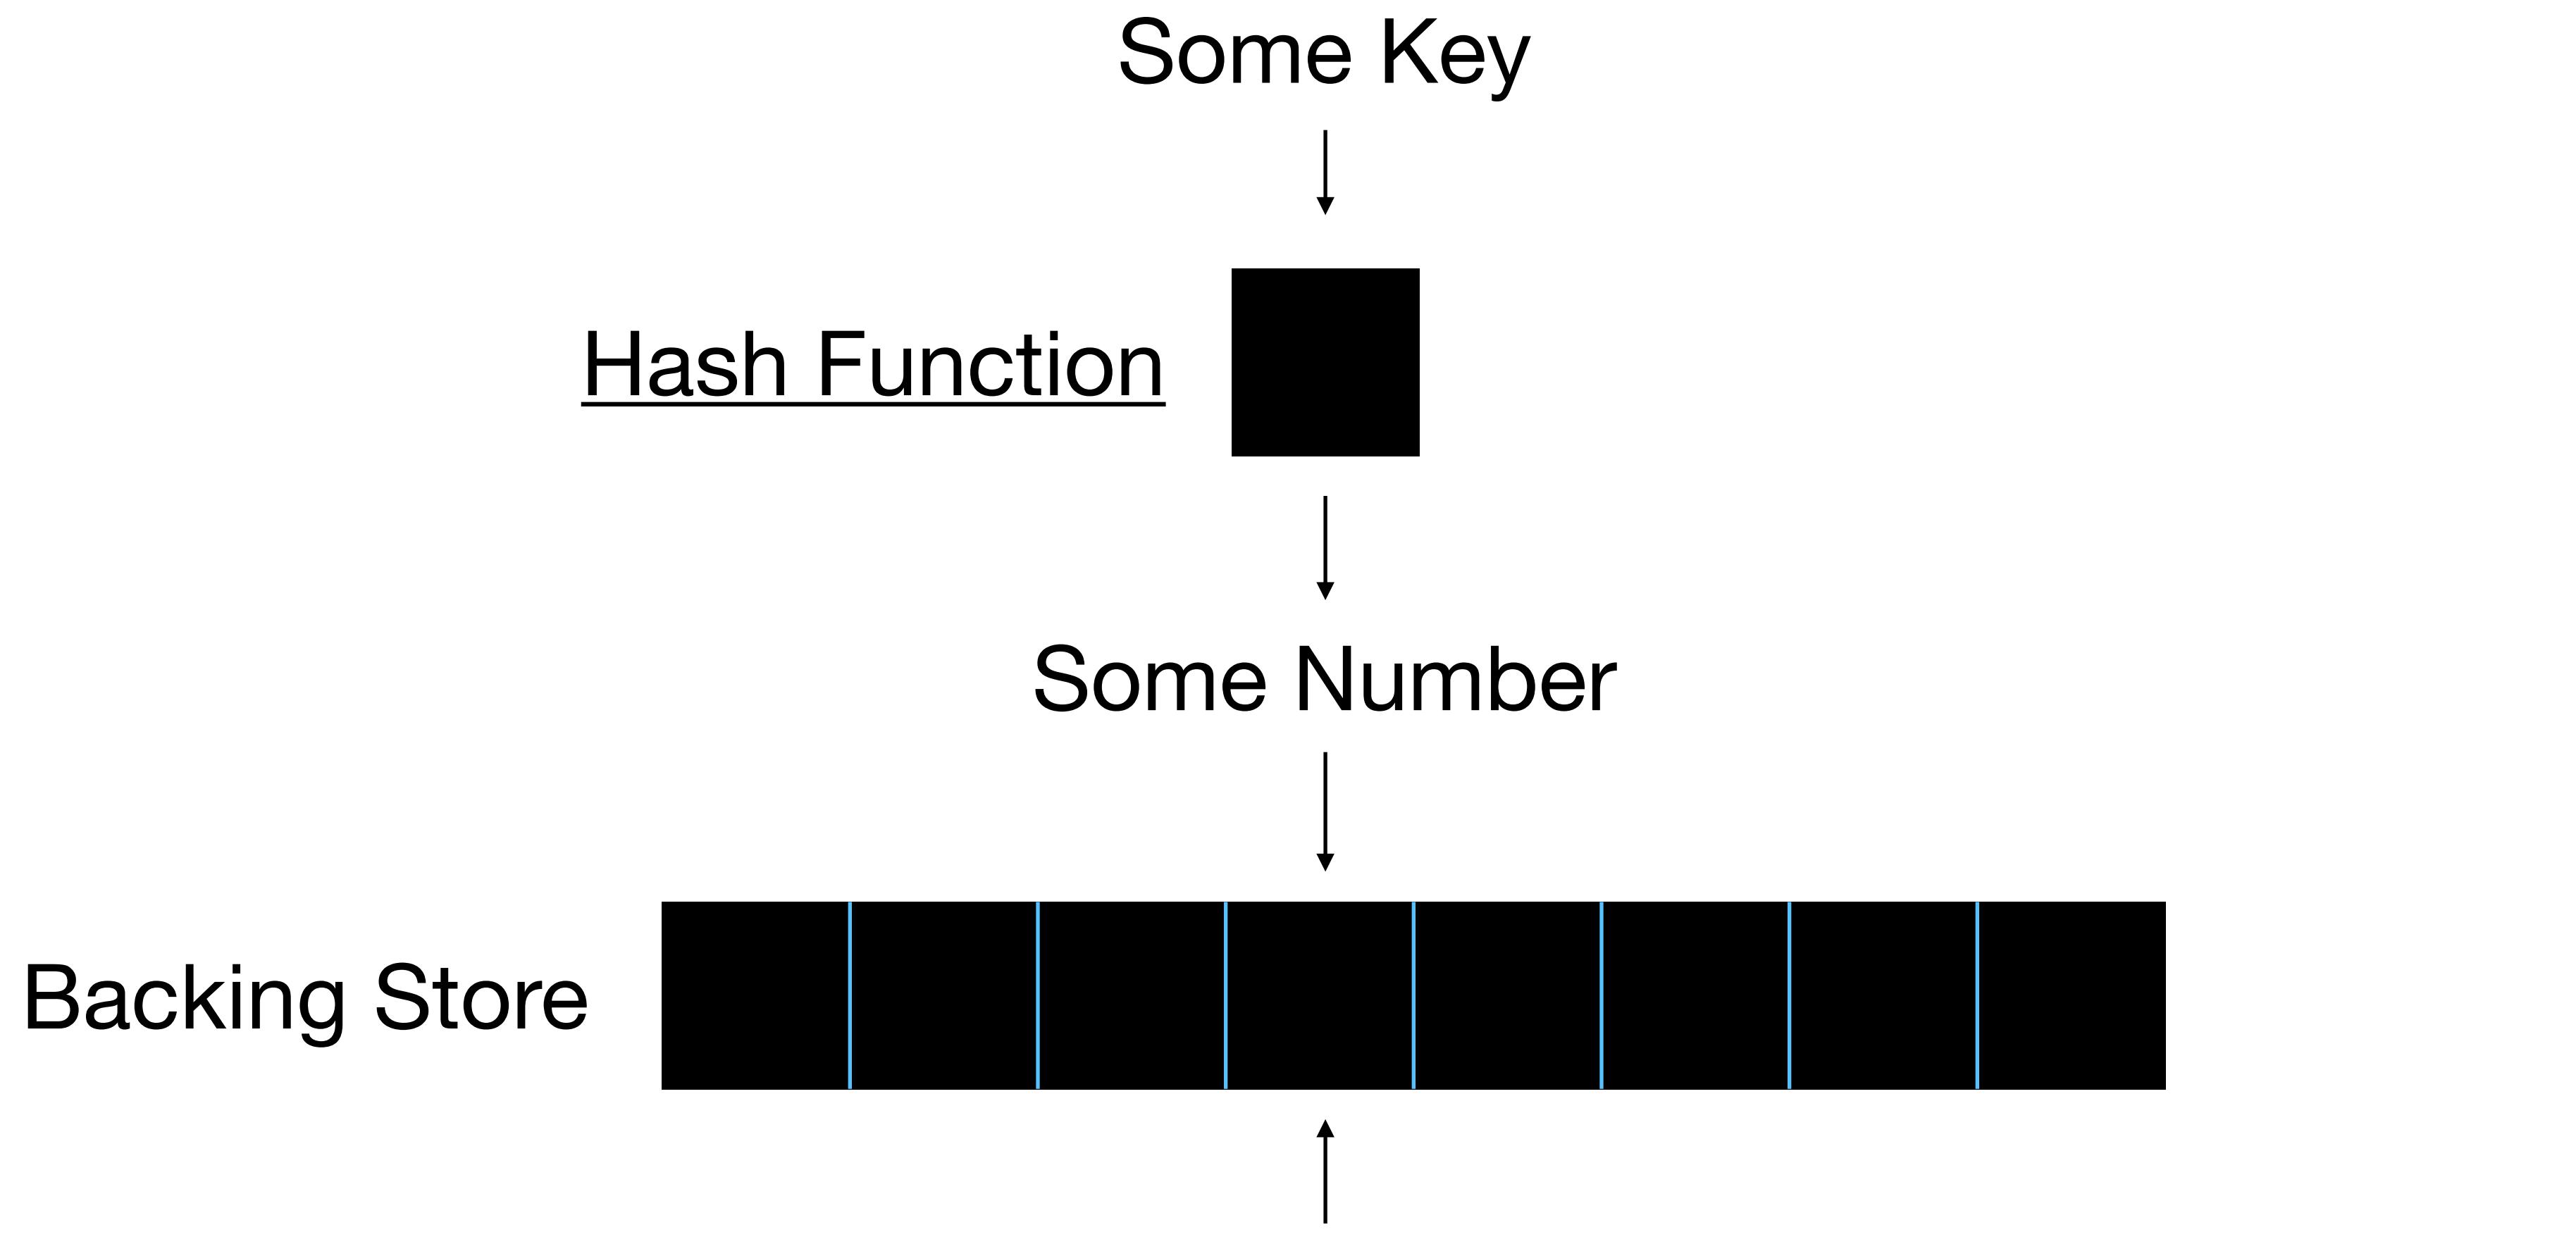
Windows is getting [dtrace support](#), so if you try this out please let us know how it goes :D

<https://github.com/flamegraph-rs/flamegraph>

Measure



HashMap Refresh

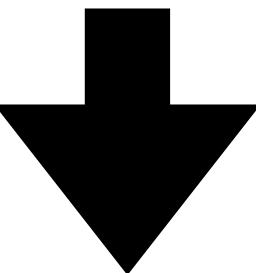


Multiple keys will map to the same array index

Collisions need to be handled and keys need to be compared

The Change

```
struct Map {  
    dog_to_food: HashMap<String, Vec<String>>,  
    food_to_ingredient: ...  
    ...  
}
```



```
struct Map {  
    dog_to_food: Vec<Vec<u16>>,  
    food_to_ingredient: ...  
    ...  
    dog_ids: Vec<String>,  
    ...  
}
```

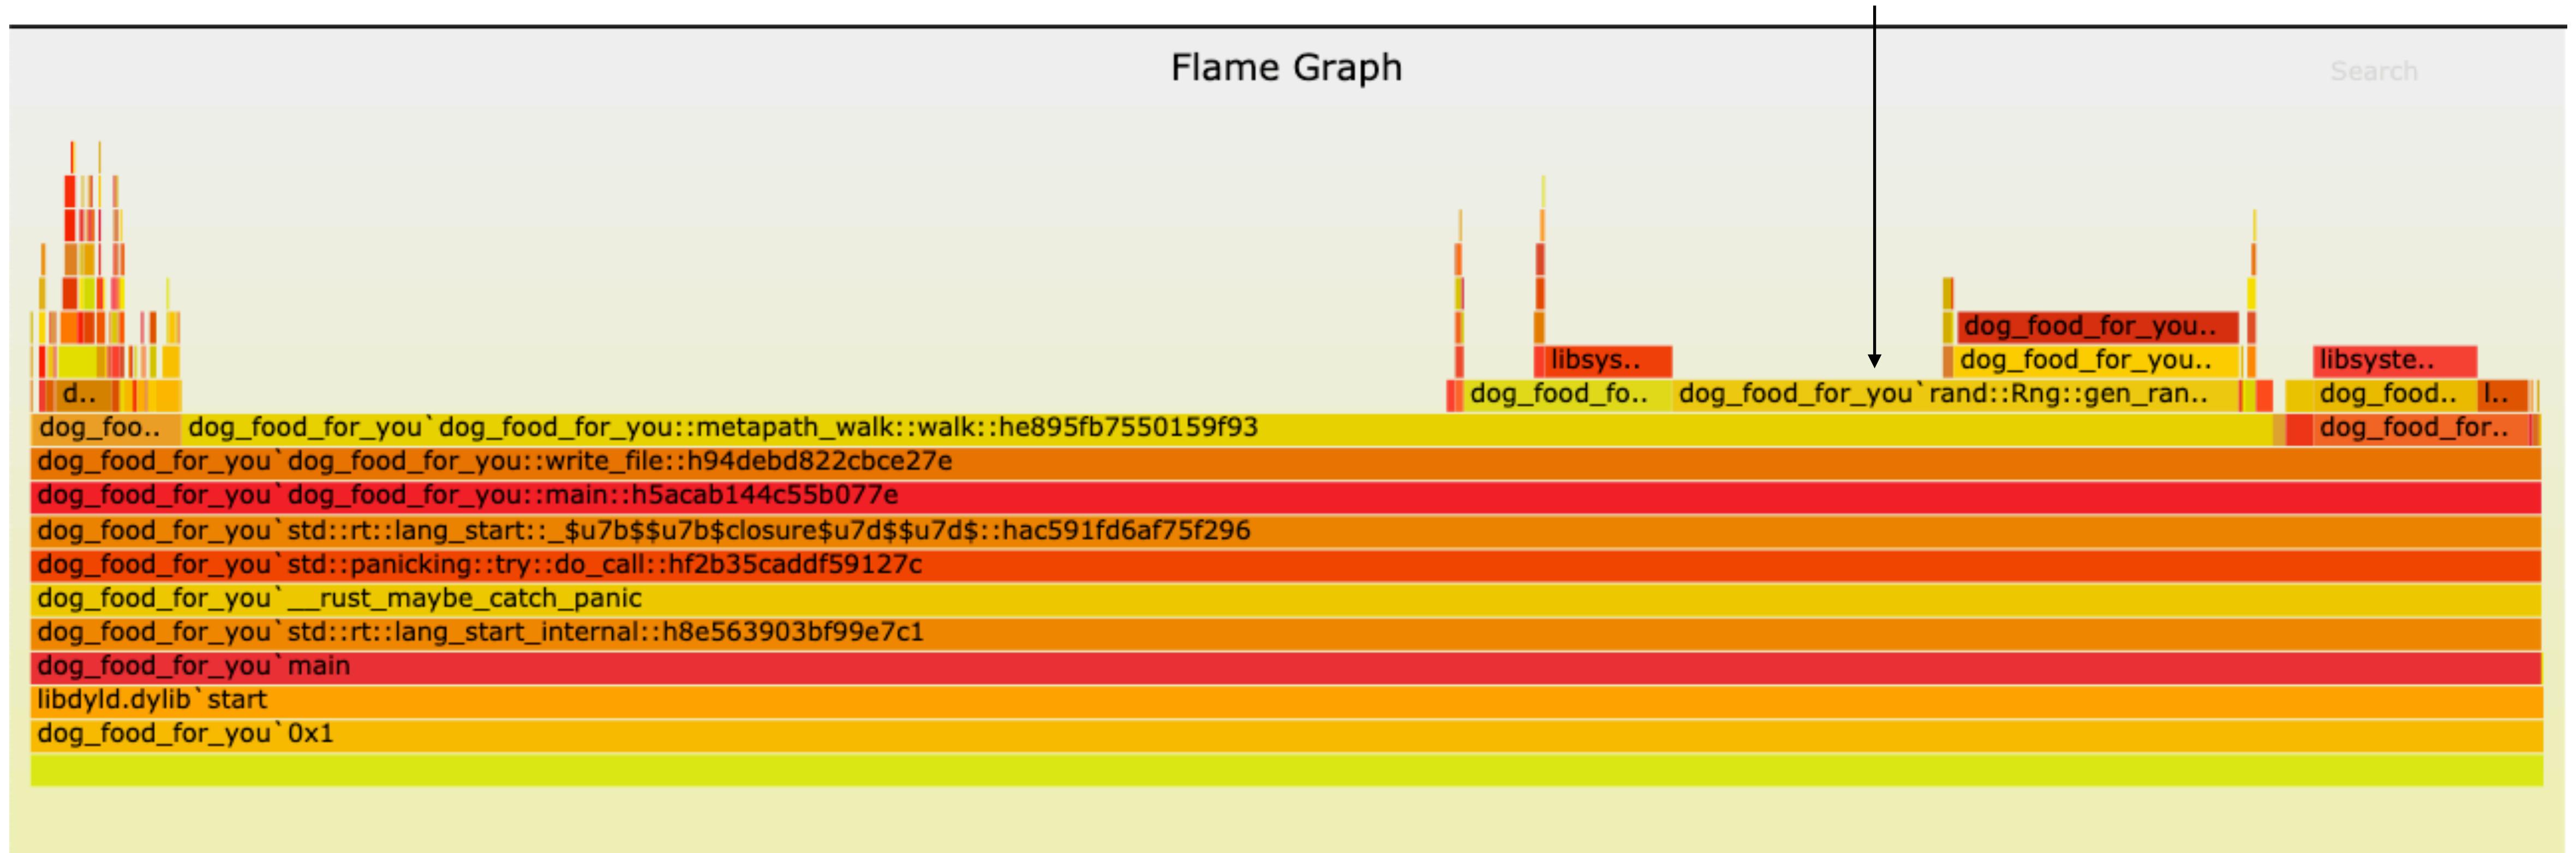
Test it out!

	Run 1 (ms)	Run 2 (ms)	Run 3 (ms)	Run 4 (ms)	Run 5 (ms)
HashMap	58291	58060	55639	55264	56900
Vec	16052	15311	15865	16001	15324
diff	-72.46%	-73.63%	-71.49%	-71.05%	-73.07%

Optimization Strategy 2.5

The Rabbit Hole

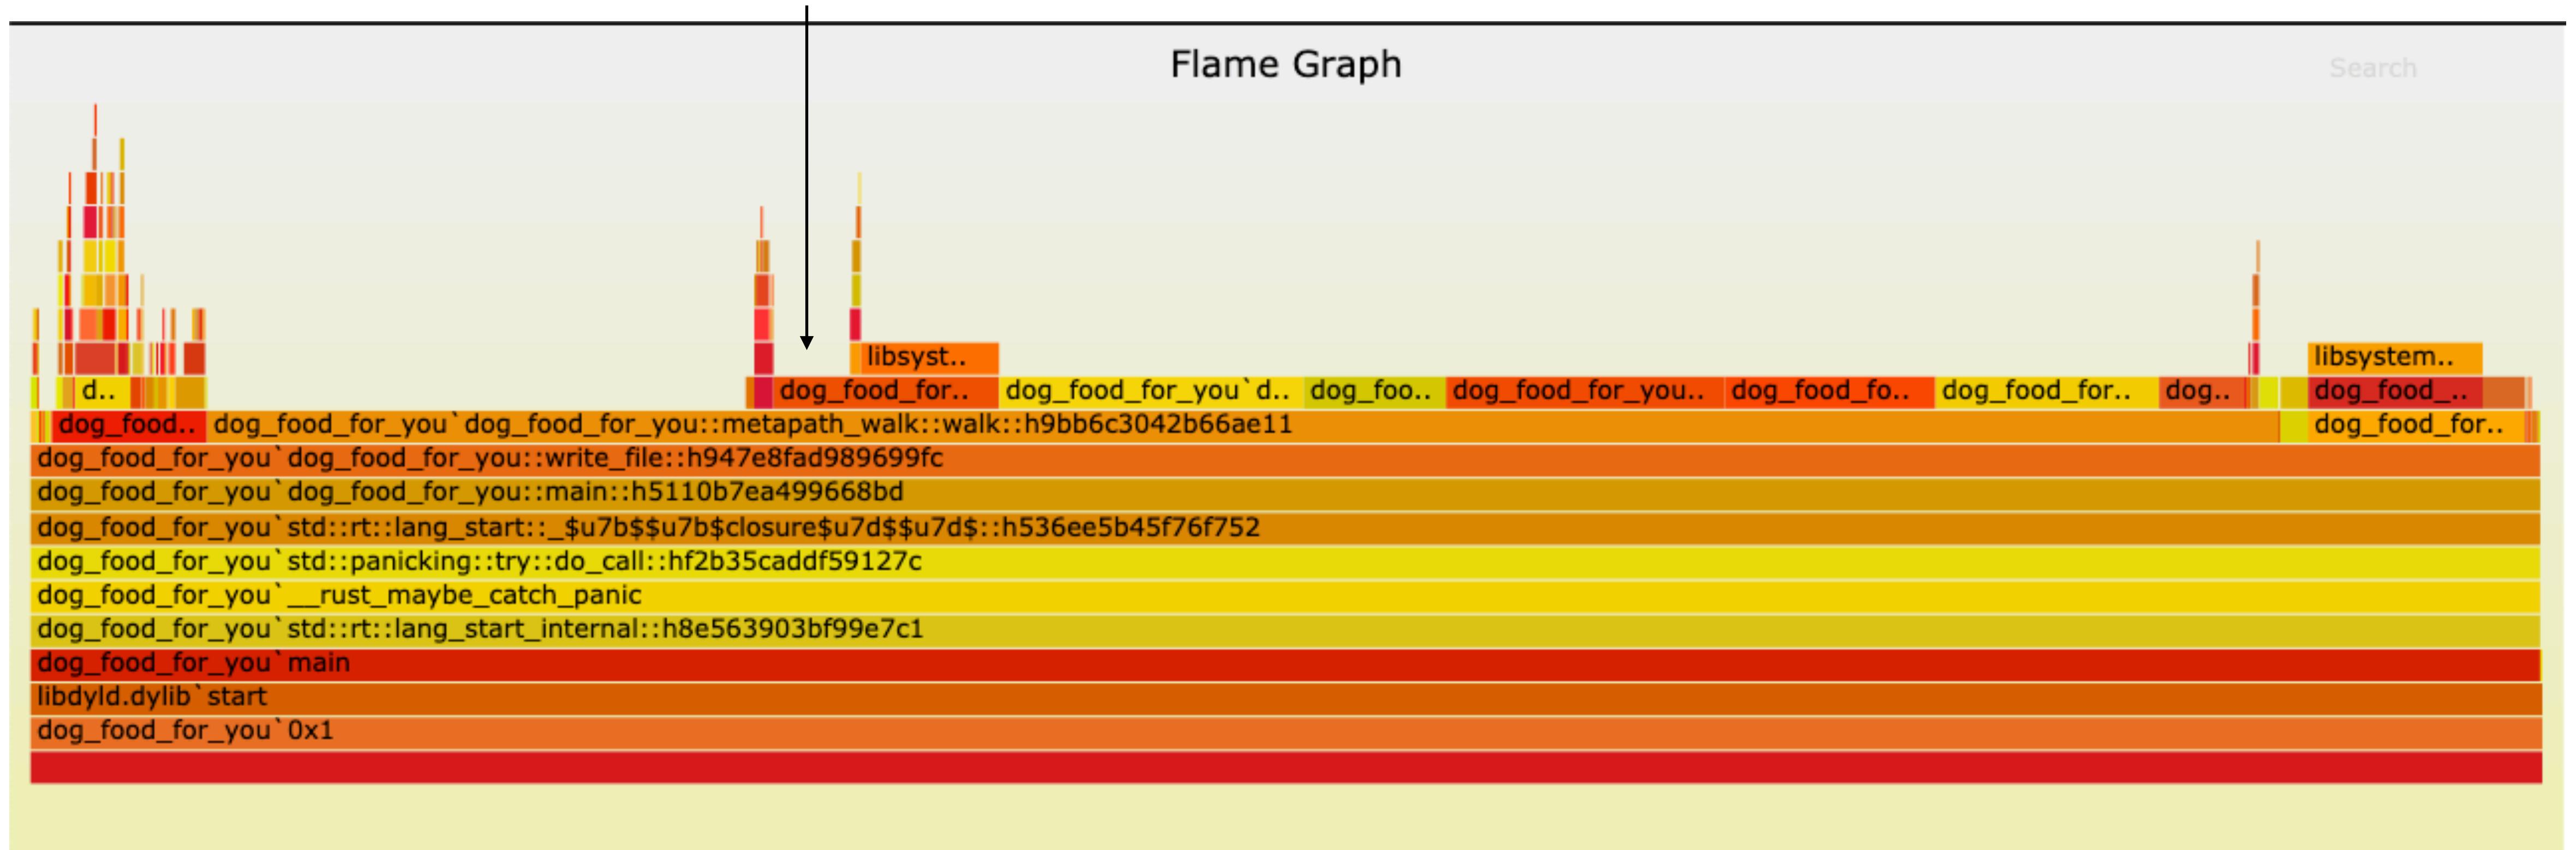
rand::Rng::gen_range 22.50%



Optimization Strategy 2.75

The Rabbit Hole

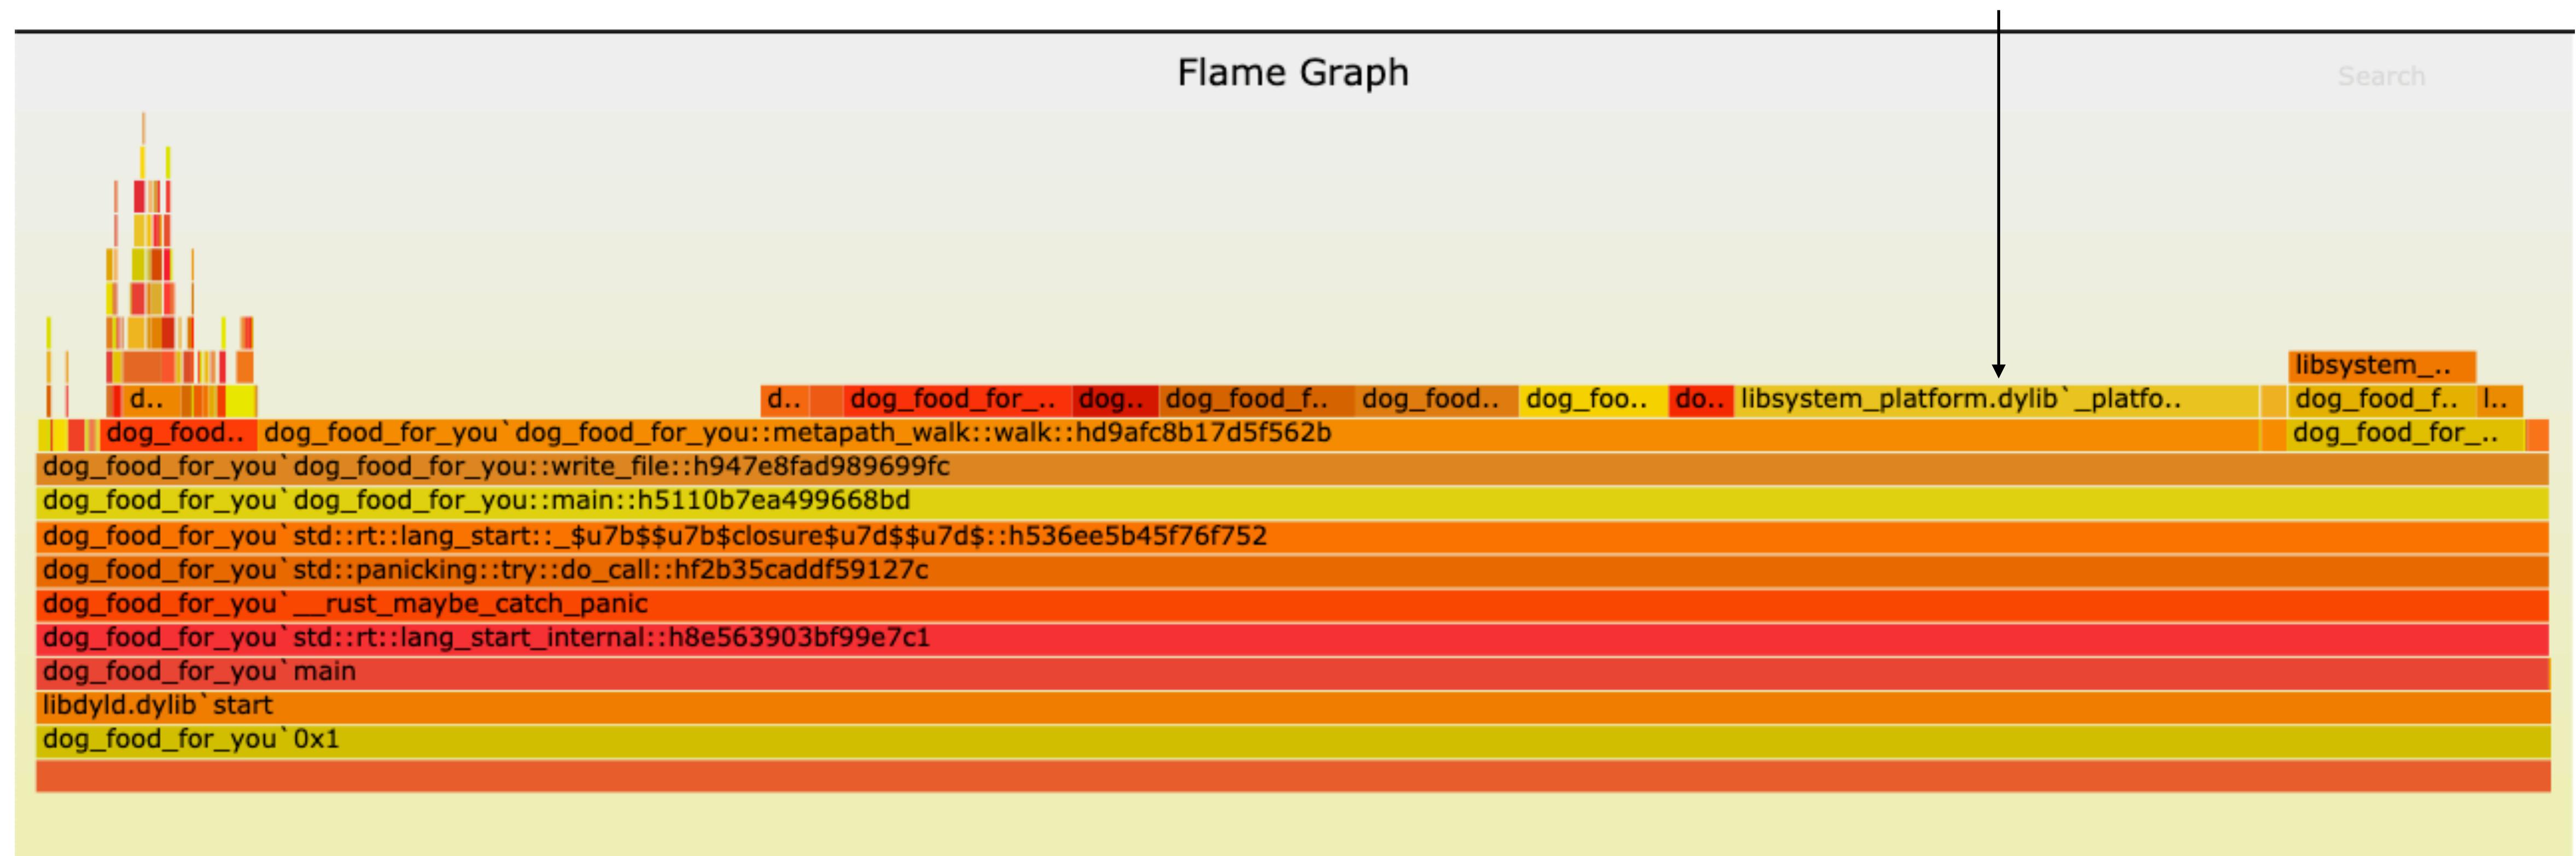
join 8.97%



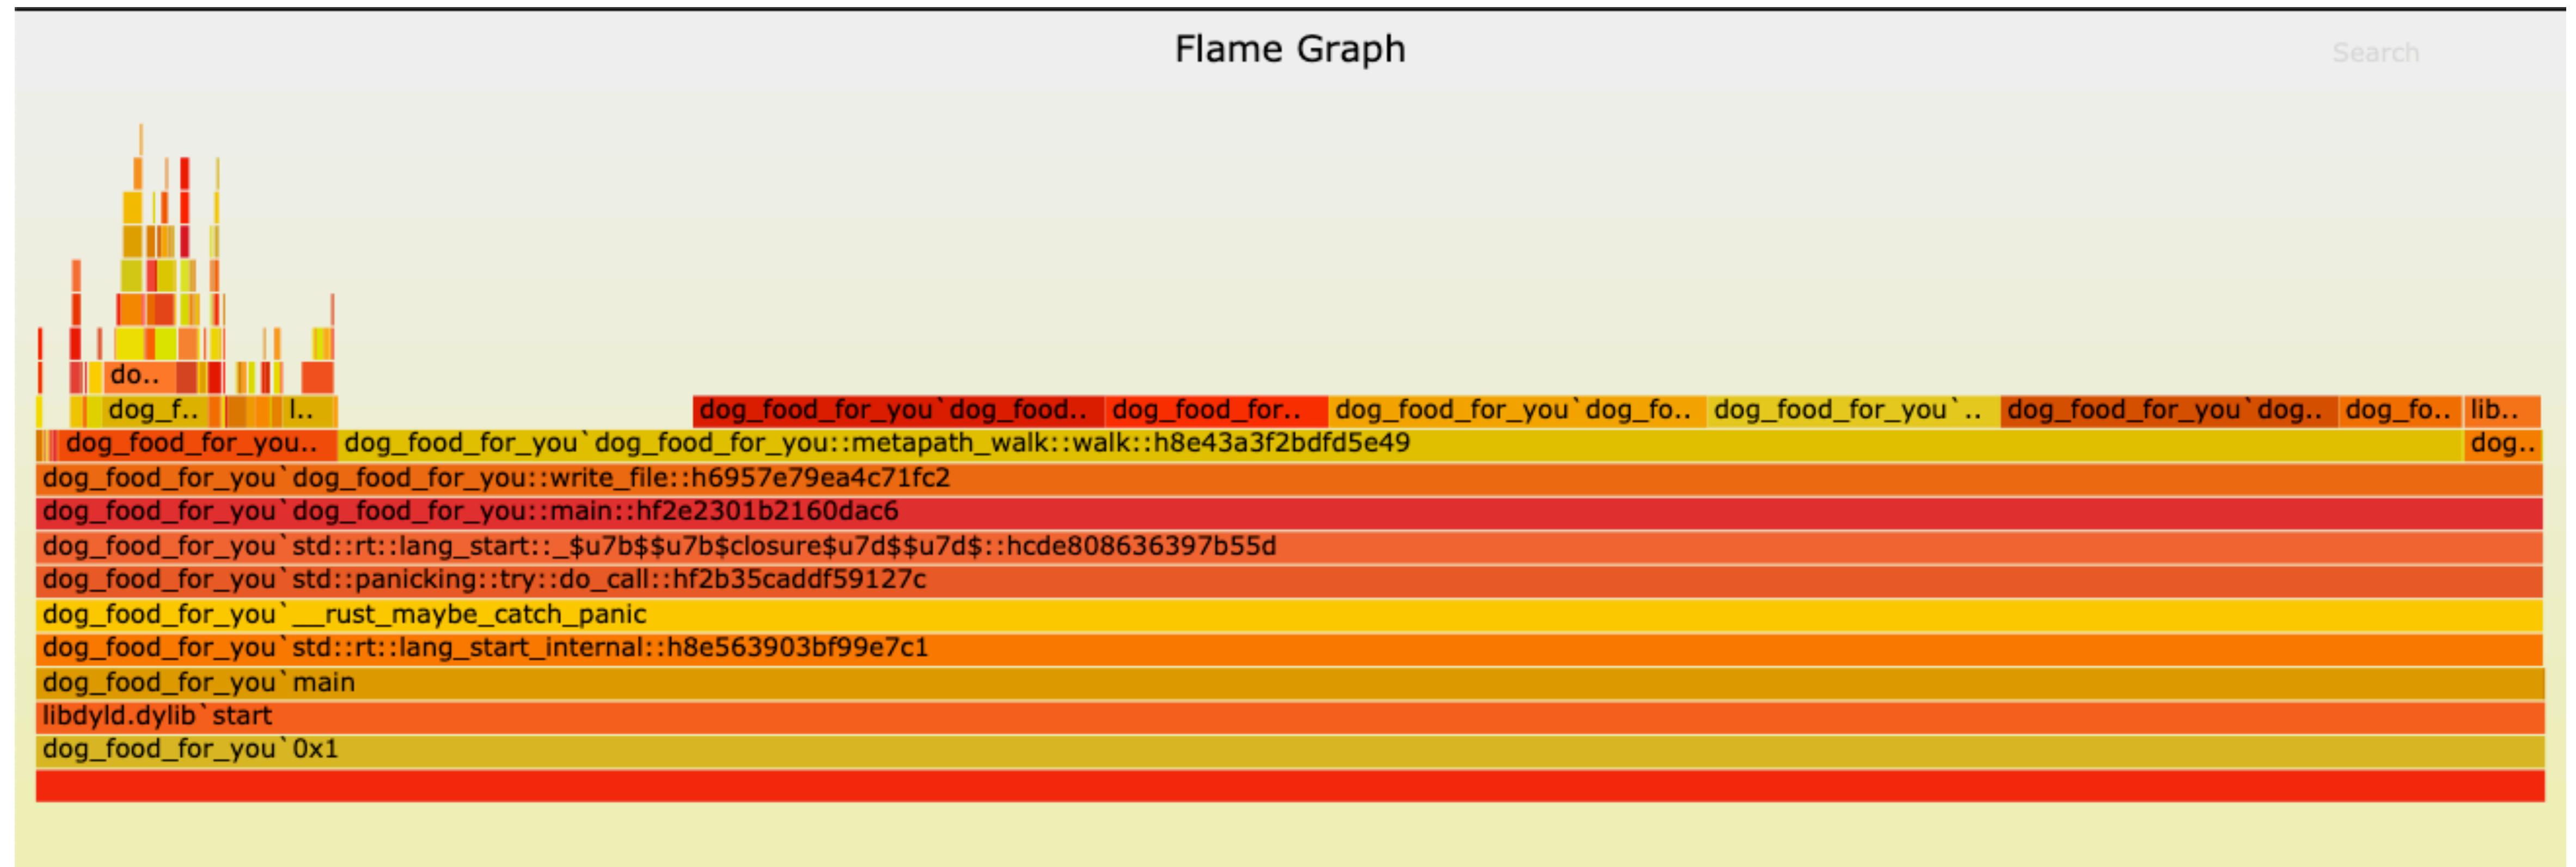
Optimization Strategy 2.875

The Rabbit Hole

memmove 20.85%



Optimization Strategy 2.9375



Test it out!

	Run 1 (ms)	Run 2 (ms)	Run 3 (ms)	Run 4 (ms)	Run 5 (ms)
Vec	18275	18592	18182	18975	18500
Bottom of the Rabbit Hole	10193	10068	9977	9864	10031
diff	-44.22%	-45.85%	-45.13%	-48.02%	-45.78%

Profiling code helps focus efforts

- Algorithmic changes often result in larger changes than micro-optimizations
- It's easy to get fall down the rabbit hole fixing less important things

Optimization Strategy 3

Fearless Concurrency



Optimization Strategy 3

Fearless Concurrency

```
pub fn write_file(...) {
    for dog in mapping.dogs() {
        for _ in 0..lines_per_dog {
            ...
        }
    }
}
```

- <https://github.com/rayon-rs/rayon>
- <https://github.com/crossbeam-rs/crossbeam>

The Diff

```
+     let (sender, receiver) = unbounded();  
+  
+     thread::spawn(move || {  
+         for line in receiver {  
+             output_file.write_all(line.as_bytes()).unwrap();  
+         }  
+         output_file.flush().unwrap();  
+     });  
  
-         for dog in mapping.dogs() {  
-             for _ in 0..lines_per_dog {  
-                 (0..lines_per_dog).into_par_iter().for_each(|_| {  
-                     let mut rng = rand::thread_rng();  
-                     let line = walk(walks_per_line, &dog, &mapping, &  
-                         output_file.write_all(line.as_bytes()).unwrap());  
-                 })  
-             }  
-             sender.clone().send(line).unwrap();  
-         }  
-     }  
- }  
+     for dog in mapping.dogs() {  
+         (0..lines_per_dog).into_par_iter().for_each(|_| {  
+             let mut rng = rand::thread_rng();  
+             let line = walk(walks_per_line, &dog, &mapping, &  
+                 output_file.write_all(line.as_bytes()).unwrap());  
+             sender.clone().send(line).unwrap();  
+         })  
+     }  
+ }  
+ }
```

Send lines to
thread for writing

Change for loop into
into_par_iter()

Test it out!

	Run 1 (ms)	Run 2 (ms)	Run 3 (ms)	Run 4 (ms)	Run 5 (ms)
Original HashMap	61582	65611	61265	61307	63206
Parallelized HashMap	11377	11355	10931	11035	11528
diff	-81.53%	-82.69%	-82.16%	-82%	-81.76%

Combining these strategies

	Run 1 (ms)	Run 2 (ms)	Run 3 (ms)	Run 4 (ms)	Run 5 (ms)
Bottom of the Rabbit Hole	7507	7650	8116	8333	7294
Parallelized	1985	1943	2125	2118	2071
diff	-73.56%	-74.6%	-73.82%	-74.58%	-71.61%

Concurrency can be an easy win

- Note - some programs can't be made concurrent!

Summary

- Guessing about performance is hard
- Flamegraphs are a great tool to inform optimization choices
- Algorithmic changes are often more impactful than micro-optimizations
- Fearless concurrency is awesome

